Execution evidence

Mainuddin Talukdar | 81255475 | Updated 22 October 2026

The original lab logs, five screenshots and fresh local CPU test transcripts accompany this index. Screenshots are genuine user-supplied lab captures. No terminal output has been fabricated.

Completed GPU experiments

Run	Evidence
3090 baseline, seed 42, LR 0.1	50 epochs; complete terminal log
3090 attention, seed 42, LR 0.1	50 epochs; original complete log recovered from Git 71fc796
3090 attention, seed 42, LR 0.001	100 epochs; complete terminal log
4070 baseline, seed 43, LR 0.001	100 epochs; complete terminal log
4070 attention, seed 43, LR 0.001	300 epochs; terminal log covers 201-300; CSV covers 1-300

Both machines use Python 3.12.3 and TensorFlow 2.21.0. The recorded devices are GPU:0 on cs23002lv (RTX 3090) and cs25049kx (RTX 4070). Configurations and runtime metadata are preserved in results/. The latest attention runtime records a resume from 200 to 300 epochs.

Validation and source provenance

The updated source passed all five assertion-based tests on the Mac CPU; the supplied print test was also executed successfully. Logs are in evidence/tests.txt and evidence/supplied-tests.txt. These are local CPU checks, not GPU training runs.

The analysis checked all 600 epoch rows, available matching log entries, recorded source hashes, and 60 saved raw/display image pairs. The raw/display conversion and recomputed image MSE agree with the recorded data. The training source differs between early and later runs only in its default learning rate. code/historical/project_main_lr01.py preserves the original source; explicit --learning-rate arguments reproduce either setting.

evidence/overwritten-attention-s42.log records a rejected attempt to reuse a non-empty run folder. It does not invalidate the completed earlier run; that run is supported by the recovered log and unchanged results. The missing early 4070 terminal segment is explicitly acknowledged rather than reconstructed.

Reproduction

Use code/run_experiments.sh with a separately supplied chest_50.pickle and a new output folder. The script runs tests, then the lower-rate seed-43 baseline for 100 epochs and attention for 300 epochs. compare_runs.py compares their common epoch-100 snapshot. README_SUBMISSION.txt explains installation and the archive contents. Training checkpoints and datasets are excluded; fresh training is reproducible from the provided source.